

# Copy\_of\_Copy\_of\_4\_DonorsChoose\_NB

March 16, 2019

## 1 DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result, there are three main p

How to scale current manual processes and resources to screen 500,000 projects so that they can be posted as quickly a

<li>How to increase the consistency of project vetting across different volunteers to improve the experience for teach

<li>How to focus volunteer time on the applications that need the most assistance</li>

</ul>

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

### 1.1 About the DonorsChoose Data Set

The train.csv data set provided by DonorsChoose contains the following features:

| Feature    | Description   |
|------------|---|
| project_id | A unique identifier for the proposed project. <b>Example:</b> p036502 |

project\_title | Title of the project. **Examples:**

Art Will Make You Happy!

First Grade Fun

project\_grade\_category | Grade level of students for which the project is targeted. One of the following enumerated values:

Grades PreK-2

Grades 3-5

Grades 6-8

Grades 9-12

project\_subject\_categories | One or more (comma-separated) subject categories for the project from the following enumerated list of values:

Applied Learning  
 Care & Hunger  
 Health & Sports  
 History & Civics  
 Literacy & Language  
 Math & Science  
 Music & The Arts  
 Special Needs  
 Warmth

**Examples:**

Music & The Arts  
 Literacy & Language, Math & Science

school\_state | State where school is located ([Two-letter U.S. postal code](#)). **Example:** WY  
 project\_subject\_subcategories | One or more (comma-separated) subject subcategories for the project. **Examples:**

Literacy  
 Literature & Writing, Social Sciences

project\_resource\_summary | An explanation of the resources needed for the project. **Example:**  
 My students need hands on literacy materials to manage sensory needs!

project\_essay\_1 | First application essay

project\_essay\_2 | *Second application essay* project\_essay\_3 | Third application essay  
 project\_essay\_4 | *Fourth application essay* project\_submitted\_datetime | Datetime when project application was submitted. **Example:** 2016-04-28 12:43:56.245

teacher\_id | A unique identifier for the teacher of the proposed project. **Example:**  
 bdf8baa8fedef6bfec7ae4ff1c15c56

teacher\_prefix | Teacher's title. One of the following enumerated values:

nan  
 Dr.  
 Mr.  
 Mrs.  
 Ms.  
 Teacher.

teacher\_number\_of\_previously\_posted\_projects | Number of project applications previously submitted by the same teacher. **Example:** 2

\* See the section Notes on the Essay Data for more details about these features.

Additionally, the resources.csv data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

| Feature     | Description   |
|-------------|---|
| id          | A project_id value from the train.csv file. <b>Example:</b><br>p036502              |
| description | Description of the resource. <b>Example:</b><br>Tenor Saxophone<br>Reeds, Box of 25 |

| Feature  | Description  |
|----------|--|
| quantity | Quantity of the resource required.                   |
| price    | Price of the resource required. <b>Example:</b> 9.95 |

**Note:** Many projects require multiple resources. The id value corresponds to a project\_id in train.csv, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

| Label               | Description  |
|---------------------|--|
| project_is_approved | Arbitrary flag indicating whether DonorsChoose approved the project. A value of 0 indicates the project was not approved, and a value of 1 indicates the project was approved. |

### 1.1.1 Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

**project\_essay\_1:** "Introduce us to your classroom"

**project\_essay\_2:** "Tell us more about your students"

**project\_essay\_3:** "Describe how your students will use the materials you're requesting"

**project\_essay\_3:** "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

**project\_essay\_1:** "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."

**project\_essay\_2:** "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with project\_submitted\_datetime of 2016-05-17 and later, the values of project\_essay\_3 and project\_essay\_4 will be NaN.

```
In [0]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
```

```

import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from plotly import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter

```

## 1.2 1.1 Reading Data

In [0]: #since im using google\_colab, i have to mount the gdrive folder for accessing the files

```

from google.colab import drive
drive.mount('/content/gdrive')

```

Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf)

Enter your authorization code:

uuuuuuuuuuuu

Mounted at /content/gdrive

In [0]: #reading the datasets, i have taken only 5000 datapoints into consideration for avoiding mermory issues

```

project_data = pd.read_csv('/content/gdrive/My Drive/Colab Notebooks/Assignments_DonorsChoose_2
resource_data = pd.read_csv('/content/gdrive/My Drive/Colab Notebooks/Assignments_DonorsChoose_2

In [0]: print("Number of data points in train data", project_data.shape)
        print('-'*50)
        print("The attributes of data :", project_data.columns.values)

```

Number of data points in train data (50000, 17)

```

-----
The attributes of data : ['Unnamed: 0' 'id' 'teacher_id' 'teacher_prefix' 'school_state'
'project_submitted_datetime' 'project_grade_category'
'project_subject_categories' 'project_subject_subcategories'
'project_title' 'project_essay_1' 'project_essay_2' 'project_essay_3'
'project_essay_4' 'project_resource_summary'
'teacher_number_of_previously_posted_projects' 'project_is_approved']

```

```

In [0]: # how to replace elements in list python: https://stackoverflow.com/a/2582163/4084039
        cols = ['Date' if x=='project_submitted_datetime' else x for x in list(project_data.columns)]

        #sort dataframe based on time pandas python: https://stackoverflow.com/a/49702492/4084039
        project_data['Date'] = pd.to_datetime(project_data['project_submitted_datetime'])
        project_data.drop('project_submitted_datetime', axis=1, inplace=True)
        project_data.sort_values(by=['Date'], inplace=True)

        # how to reorder columns pandas python: https://stackoverflow.com/a/13148611/4084039
        project_data = project_data[cols]

```

```

project_data.head(2)

Out[0]:      Unnamed: 0      id      teacher_id teacher_prefix \
473      100660 p234804 cbc0e38f522143b86d372f8b43d4cff3      Mrs.
41558      33679 p137682 06f6e62e17de34fcf81020c77549e1d5      Mrs.

      school_state      Date project_grade_category \
473      GA 2016-04-27 00:53:00      Grades PreK-2
41558      WA 2016-04-27 01:05:25      Grades 3-5

      project_subject_categories project_subject_subcategories \
473      Applied Learning      Early Development
41558      Literacy & Language      Literacy

      project_title \
473      Flexible Seating for Flexible Learning
41558      Going Deep: The Art of Inner Thinking!

```

```

                                project_essay_1 \
473  I recently read an article about giving studen...
41558 My students crave challenge, they eat obstacle...

```

```

                                project_essay_2 \
473  I teach at a low-income (Title 1) school. Ever...
41558 We are an urban, public k-5 elementary school...

```

```

                                project_essay_3 \
473  We need a classroom rug that we can use as a c...
41558 With the new common core standards that have b...

```

```

                                project_essay_4 \
473  Benjamin Franklin once said, \"Tell me and I f...
41558 These remarkable gifts will provide students w...

```

```

                                project_resource_summary \
473  My students need flexible seating in the class...
41558 My students need copies of the New York Times ...

```

```

teacher_number_of_previously_posted_projects project_is_approved
473                                     2                 1
41558                                2                 1

```

```

In [0]: print("Number of data points in train data", resource_data.shape)
        print(resource_data.columns.values)
        resource_data.head(2)

```

Number of data points in train data (1541272, 4)  
['id' 'description' 'quantity' 'price']

```

Out[0]:      id      description  quantity \
0  p233245  LC652 - Lakeshore Double-Space Mobile Drying Rack      1
1  p069063      Bouncy Bands for Desks (Blue support pipes)      3

      price
0  149.00
1   14.95

```

### 1.3 1.2 preprocessing of project\_subject\_categories

```

In [0]: categories = list(project_data['project_subject_categories'].values)
        # remove special characters from list of strings python: https://stackoverflow.com/a/47301924/4084039

        # https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
        # https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
        # https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

```

```

cat_list = []
for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"=> "Math
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ''(i.e removing '
        j = j.replace(' ','') # we are placing all the ' '(space) with ''(empty) ex:"Math & Science"=>"Math&
        temp+=j.strip()+" " # " abc ".strip() will return "abc", remove the trailing spaces
        temp = temp.replace('&','_') # we are replacing the & value into
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))

```

## 1.4 1.3 preprocessing of project\_subject\_subcategories

```

In [0]: sub_categories = list(project_data['project_subject_subcategories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/47301924/4084039

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

sub_cat_list = []
for i in sub_categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"=> "Math
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ''(i.e removing '
        j = j.replace(' ','') # we are placing all the ' '(space) with ''(empty) ex:"Math & Science"=>"Math&
        temp +=j.strip()+" " # " abc ".strip() will return "abc", remove the trailing spaces
        temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/4084039

```

```

my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

## 1.5 1.3 Text preprocessing

```

In [0]: # merge two column text dataframe:
project_data["essay"] = project_data["project_essay_1"].map(str) + \
    project_data["project_essay_2"].map(str) + \
    project_data["project_essay_3"].map(str) + \
    project_data["project_essay_4"].map(str)

In [0]: project_data.head(2)

Out[0]:      Unnamed: 0      id      teacher_id teacher_prefix \
473      100660 p234804 cbc0e38f522143b86d372f8b43d4cff3      Mrs.
41558      33679 p137682 06f6e62e17de34fcf81020c77549e1d5      Mrs.

      school_state      Date project_grade_category \
473      GA 2016-04-27 00:53:00      Grades PreK-2
41558      WA 2016-04-27 01:05:25      Grades 3-5

      project_title \
473  Flexible Seating for Flexible Learning
41558  Going Deep: The Art of Inner Thinking!

      project_essay_1 \
473  I recently read an article about giving studen...
41558  My students crave challenge, they eat obstacle...

      project_essay_2 \
473  I teach at a low-income (Title 1) school. Ever...
41558  We are an urban, public k-5 elementary school...

      project_essay_3 \
473  We need a classroom rug that we can use as a c...
41558  With the new common core standards that have b...

      project_essay_4 \
473  Benjamin Franklin once said, \"Tell me and I f...
41558  These remarkable gifts will provide students w...

      project_resource_summary \
473  My students need flexible seating in the class...
41558  My students need copies of the New York Times ...

```



```

teacher_number_of_previously_posted_projects project_is_approved \
473      2      1
41558    2      1

```

```

clean_categories clean_subcategories \
473    AppliedLearning    EarlyDevelopment
41558  Literacy_Language    Literacy

```

```

essay
473    I recently read an article about giving studen...
41558  My students crave challenge, they eat obstacle...

```

In [0]: ##### 1.4.2.3 Using Pretrained Models: TFIDF weighted W2V

```

In [0]: # printing some random reviews
print(project_data['essay'].values[0])
print("="*50)
print(project_data['essay'].values[150])
print("="*50)
print(project_data['essay'].values[1000])
print("="*50)
print(project_data['essay'].values[20000])
print("="*50)
#print(project_data['essay'].values[99999])
#print("="*50)

```

```

I recently read an article about giving students a choice about how they learn. We already set goals; why not let t
=====
At the beginning of every class we start out with a Math Application problem to help students see the relevance o
=====
My students love coming to school and they love learning. I strive daily to make our classroom a relaxed, comfort
=====
I teach at a Title 1 school, with 73% of my students who receive free/reduced lunch. Our school provides free bre
=====

```

In [0]: # <https://stackoverflow.com/a/47091490/4084039>

```

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)

```

```

phrase = re.sub(r"\d", " would", phrase)
phrase = re.sub(r"\ll", " will", phrase)
phrase = re.sub(r"\t", " not", phrase)
phrase = re.sub(r"\ve", " have", phrase)
phrase = re.sub(r"\m", " am", phrase)
return phrase

```

```

In [0]: sent = decontracted(project_data['essay'].values[20000])
        print(sent)
        print("="*50)

```

I teach at a Title 1 school, with 73% of my students who receive free/reduced lunch. Our school provides free bre  
=====

```

In [0]: # \r \n \t remove from string python: http://texthandler.com/info/remove-line-breaks-python/
        sent = sent.replace('\r', ' ')
        sent = sent.replace('\n', ' ')
        sent = sent.replace('\t', ' ')
        print(sent)

```

I teach at a Title 1 school, with 73% of my students who receive free/reduced lunch. Our school provides free bre

```

In [0]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
        sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
        print(sent)

```

I teach at a Title 1 school with 73 of my students who receive free reduced lunch Our school provides free breakfa

```

In [0]: # https://gist.github.com/sebleier/554280
        # we are removing the words from the stop words list: 'no', 'nor', 'not'
        stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", \
                    "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
                    'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', \
                    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', \
                    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', \
                    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', \
                    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after', \
                    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further', \
                    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', \
                    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', \
                    's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', \
                    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', \
                    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn', \
                    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', \
                    'won', "won't", 'wouldn', "wouldn't"]

```

```
In [0]: # Combining all the above stundents
from tqdm import tqdm
preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(project_data['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\\"', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
    preprocessed_essays.append(sent.lower().strip())
```

100%| 50000/50000 [00:32<00:00, 1552.22it/s]

```
In [0]: # after preprocessing
preprocessed_essays[20000]
```

Out[0]: 'teach title 1 school 73 students receive free reduced lunch school provides free breakfast students special c

## 1.4 Preprocessing of project\_title

```
In [0]: # similarly you can preprocess the titles also
```

```
from tqdm import tqdm
preprocessed_project_title = []
# tqdm is for printing the status bar
for sentence in tqdm(project_data['project_title'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\\"', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e not in stopwords)
    preprocessed_project_title.append(sent.lower().strip())
```

100%| 50000/50000 [00:01<00:00, 32611.76it/s]

## 1.6 1.5 Preparing data for models

```
In [0]: project_data.columns
```

Out[0]: Index(['Unnamed: 0', 'id', 'teacher\_id', 'teacher\_prefix', 'school\_state',  
'Date', 'project\_grade\_category', 'project\_title', 'project\_essay\_1',  
'project\_essay\_2', 'project\_essay\_3', 'project\_essay\_4',

```

        'project_resource_summary',
        'teacher_number_of_previously_posted_projects', 'project_is_approved',
        'clean_categories', 'clean_subcategories', 'essay'],
        dtype='object')

```

we are going to consider

- school\_state : categorical data
- clean\_categories : categorical data
- clean\_subcategories : categorical data
- project\_grade\_category : categorical data
- teacher\_prefix : categorical data
  
- project\_title : text data
- text : text data
- project\_resource\_summary: text data (optinal)
  
- quantity : numerical (optinal)
- teacher\_number\_of\_previously\_posted\_projects : numerical
- price : numerical

### 1.6.1 Modifying DataSet (essay & project\_title)

```

In [0]: project_data['clean_essay'] = preprocessed_essays
        project_data['clean_project_title'] = preprocessed_project_title
        project_data.drop(['essay'], axis=1, inplace=True)
        project_data.drop(['project_title'], axis=1, inplace=True)

```

```

In [0]: project_data.head(1)

```

```

Out[0]:   Unnamed: 0      id      teacher_id teacher_prefix \
473      100660  p234804  cbc0e38f522143b86d372f8b43d4cff3  Mrs.

      school_state      Date project_grade_category \
473      GA 2016-04-27 00:53:00      Grades PreK-2

      project_essay_1 \
473  I recently read an article about giving studen...

      project_essay_2 \
473  I teach at a low-income (Title 1) school. Ever...

      project_essay_3 \
473  We need a classroom rug that we can use as a c...

      project_essay_4 \
473  Benjamin Franklin once said, \"Tell me and I f...

      project_resource_summary \

```

```

473 My students need flexible seating in the class...

      teacher_number_of_previously_posted_projects project_is_approved \
473                                     2                1

      clean_categories clean_subcategories \
473 AppliedLearning EarlyDevelopment

      clean_essay \
473 recently read article giving students choice l...

      clean_project_title
473 flexible seating flexible learning

```

## 1.7 Splitting DataSet

```

In [0]: y = project_data['project_is_approved'].values
        project_data.drop(['project_is_approved'], axis=1, inplace=True)
        project_data.head(1)

Out[0]: Unnamed: 0      id      teacher_id teacher_prefix \
473      100660 p234804 cbc0e38f522143b86d372f8b43d4cff3      Mrs.

      school_state      Date project_grade_category \
473      GA 2016-04-27 00:53:00      Grades PreK-2

      project_essay_1 \
473 I recently read an article about giving studen...

      project_essay_2 \
473 I teach at a low-income (Title 1) school. Ever...

      project_essay_3 \
473 We need a classroom rug that we can use as a c...

      project_essay_4 \
473 Benjamin Franklin once said, \"Tell me and I f...

      project_resource_summary \
473 My students need flexible seating in the class...

      teacher_number_of_previously_posted_projects clean_categories \
473                                     2 AppliedLearning

      clean_subcategories      clean_essay \
473 EarlyDevelopment recently read article giving students choice l...

      clean_project_title

```

```
In [0]: X = project_data
```

```
In [0]: # please write all the code with proper documentation, and proper titles for each subsection
# go through documentations and blogs before you start coding
# first figure out what to do, and then think about how to do.
# reading and understanding error messages will be very much helpfull in debugging your code
# when you plot any graph make sure you use
# a. Title, that describes your plot, this will be very helpful to the reader
# b. Legends if needed
# c. X-axis label
# d. Y-axis label

# train test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, stratify=y)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.33, stratify=y_train)
```

### 1.7.1 1.5.1 Vectorizing Categorical data

- <https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

### 1.7.2 clean\_categories

```
In [0]: vectorizer = CountVectorizer()
vectorizer.fit(X_train['clean_categories'].values) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_cc_ohe = vectorizer.transform(X_train['clean_categories'].values)
X_cv_cc_ohe = vectorizer.transform(X_cv['clean_categories'].values)
X_test_cc_ohe = vectorizer.transform(X_test['clean_categories'].values)

feat_cc = vectorizer.get_feature_names()

print("After vectorizations")
print(X_train_cc_ohe.shape, y_train.shape)
print(X_cv_cc_ohe.shape, y_cv.shape)
print(X_test_cc_ohe.shape, y_test.shape)
print(vectorizer.get_feature_names())
print("="*100)
```

After vectorizations

(22445, 9) (22445,)

(11055, 9) (11055,)

(16500, 9) (16500,)

['appliedlearning', 'care\_hunger', 'health\_sports', 'history\_civics', 'literacy\_language', 'math\_science', 'music\_a

### 1.7.3 clean\_subcategories

```
=====

In [0]: #vectorizer = CountVectorizer()
        vectorizer.fit(X_train['clean_subcategories'].values) # fit has to happen only on train data

        # we use the fitted CountVectorizer to convert the text to vector
        X_train_csc_ohe = vectorizer.transform(X_train['clean_subcategories'].values)
        X_cv_csc_ohe = vectorizer.transform(X_cv['clean_subcategories'].values)
        X_test_csc_ohe = vectorizer.transform(X_test['clean_subcategories'].values)

        feat_csc = vectorizer.get_feature_names()

        print("After vectorizations")
        print(X_train_csc_ohe.shape, y_train.shape)
        print(X_cv_csc_ohe.shape, y_cv.shape)
        print(X_test_csc_ohe.shape, y_test.shape)
        print(vectorizer.get_feature_names())
        print("="*100)
```

After vectorizations

(22445, 30) (22445,)

(11055, 30) (11055,)

(16500, 30) (16500,)

['appliedsciences', 'care\_hunger', 'charactereducation', 'civics\_government', 'college\_careerprep', 'communityserv

=====

### 1.7.4 school\_state

```
In [0]: #vectorizer = CountVectorizer()
        vectorizer.fit(X_train['school_state'].values) # fit has to happen only on train data

        # we use the fitted CountVectorizer to convert the text to vector
        X_train_state_ohe = vectorizer.transform(X_train['school_state'].values)
        X_cv_state_ohe = vectorizer.transform(X_cv['school_state'].values)
        X_test_state_ohe = vectorizer.transform(X_test['school_state'].values)

        feat_ss = vectorizer.get_feature_names()

        print("After vectorizations")
        print(X_train_state_ohe.shape, y_train.shape)
        print(X_cv_state_ohe.shape, y_cv.shape)
        print(X_test_state_ohe.shape, y_test.shape)
```

```
print(vectorizer.get_feature_names())
print("="*100)
```

After vectorizations

(22445, 51) (22445,)

(11055, 51) (11055,)

(16500, 51) (16500,)

['ak', 'al', 'ar', 'az', 'ca', 'co', 'ct', 'dc', 'de', 'fl', 'ga', 'hi', 'ia', 'id', 'il', 'in', 'ks', 'ky', 'la', 'ma', 'md', 'me', 'mi', 'm

=====

### 1.7.5 teacher\_prefix

```
In [0]: #vectorizer = CountVectorizer()
vectorizer.fit(X_train['teacher_prefix'].values.astype('U')) # fit has to happen only on train data
```

```
# we use the fitted CountVectorizer to convert the text to vector
```

```
X_train_teacher_ohe = vectorizer.transform(X_train['teacher_prefix'].values.astype('U'))
```

```
X_cv_teacher_ohe = vectorizer.transform(X_cv['teacher_prefix'].values.astype('U'))
```

```
X_test_teacher_ohe = vectorizer.transform(X_test['teacher_prefix'].values.astype('U'))
```

```
feat_tp = vectorizer.get_feature_names()
```

```
print("After vectorizations")
```

```
print(X_train_teacher_ohe.shape, y_train.shape)
```

```
print(X_cv_teacher_ohe.shape, y_cv.shape)
```

```
print(X_test_teacher_ohe.shape, y_test.shape)
```

```
print(vectorizer.get_feature_names())
```

```
print("="*100)
```

After vectorizations

(22445, 5) (22445,)

(11055, 5) (11055,)

(16500, 5) (16500,)

['dr', 'mr', 'mrs', 'ms', 'teacher']

=====

### 1.7.6 project\_grade\_category

```
In [0]: #vectorizer = CountVectorizer()
vectorizer.fit(X_train['project_grade_category'].values) # fit has to happen only on train data
```

```
# we use the fitted CountVectorizer to convert the text to vector
```

```
X_train_grade_ohe = vectorizer.transform(X_train['project_grade_category'].values)
```

```
X_cv_grade_ohe = vectorizer.transform(X_cv['project_grade_category'].values)
```

```
X_test_grade_ohe = vectorizer.transform(X_test['project_grade_category'].values)
```



```
feat_pgc = vectorizer.get_feature_names()
```

```
print("After vectorizations")
print(X_train_grade_ohe.shape, y_train.shape)
print(X_cv_grade_ohe.shape, y_cv.shape)
print(X_test_grade_ohe.shape, y_test.shape)
print(vectorizer.get_feature_names())
print("="*100)
```

After vectorizations

(22445, 3) (22445,)

(11055, 3) (11055,)

(16500, 3) (16500,)

['12', 'grades', 'prek']

=====

## 1.7.7 1.5.2 Vectorizing Text data

### 1.5.2.1 Bag of words

#### 1.7.8 essays

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
vectorizer.fit(X_train['clean_essay'].values) # fit has to happen only on train data
feat_bow_e = vectorizer.get_feature_names()
```

```
# we use the fitted CountVectorizer to convert the text to vector
X_train_essay_bow = vectorizer.transform(X_train['clean_essay'].values)
X_cv_essay_bow = vectorizer.transform(X_cv['clean_essay'].values)
X_test_essay_bow = vectorizer.transform(X_test['clean_essay'].values)
feat_bow = vectorizer.get_feature_names()
```

```
print("After vectorizations")
print(X_train_essay_bow.shape, y_train.shape)
print(X_cv_essay_bow.shape, y_cv.shape)
print(X_test_essay_bow.shape, y_test.shape)
print("="*100)
```

After vectorizations

(22445, 50000) (22445,)

(11055, 50000) (11055,)

(16500, 50000) (16500,)

=====

### 1.7.9 project\_title

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
        #vectorizer = CountVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
        vectorizer.fit(X_train['clean_project_title'].values) # fit has to happen only on train data
        feat_bow_pt = vectorizer.get_feature_names()

        # we use the fitted CountVectorizer to convert the text to vector
        X_train_pt_bow = vectorizer.transform(X_train['clean_project_title'].values)
        X_cv_pt_bow = vectorizer.transform(X_cv['clean_project_title'].values)
        X_test_pt_bow = vectorizer.transform(X_test['clean_project_title'].values)

        print("After vectorizations")
        print(X_train_pt_bow.shape, y_train.shape)
        print(X_cv_pt_bow.shape, y_cv.shape)
        print(X_test_pt_bow.shape, y_test.shape)
        print("=="*100)
```

After vectorizations  
(22445, 2010) (22445,)  
(11055, 2010) (11055,)  
(16500, 2010) (16500,)  
=====

### 1.5.2.2 TFIDF vectorizer

#### 1.7.10 essays

```
In [0]: from sklearn.feature_extraction.text import TfidfVectorizer
        vectorizert = TfidfVectorizer(min_df=10)
        X_train_essay_tfidf = vectorizert.fit_transform(X_train['clean_essay'].values)
        X_cv_essay_tfidf = vectorizert.transform(X_cv['clean_essay'].values)
        X_test_essay_tfidf = vectorizert.transform(X_test['clean_essay'].values)

        feat_tfidf_e = vectorizert.get_feature_names()

        print("Shape of matrix after one hot encodig ",X_train_essay_tfidf.shape)
        print("Shape of matrix after one hot encodig ",X_cv_essay_tfidf.shape)
        print("Shape of matrix after one hot encodig ",X_test_essay_tfidf.shape)
```

Shape of matrix after one hot encodig (22445, 8786)  
Shape of matrix after one hot encodig (11055, 8786)  
Shape of matrix after one hot encodig (16500, 8786)

### 1.7.11 project\_title

```
In [0]: from sklearn.feature_extraction.text import TfidfVectorizer
        #vectorizer = TfidfVectorizer(min_df=10)
        X_train_pt_tfidf = vectorizer.fit_transform(X_train['clean_project_title'].values)
        X_cv_pt_tfidf = vectorizer.transform(X_cv['clean_project_title'].values)
        X_test_pt_tfidf = vectorizer.transform(X_test['clean_project_title'].values)

        feat_tfidf_pt = vectorizer.get_feature_names()

        print("Shape of matrix after one hot encoding ",X_train_pt_tfidf.shape)
        print("Shape of matrix after one hot encoding ",X_cv_pt_tfidf.shape)
        print("Shape of matrix after one hot encoding ",X_test_pt_tfidf.shape)
```

Shape of matrix after one hot encoding (22445, 1227)

Shape of matrix after one hot encoding (11055, 1227)

Shape of matrix after one hot encoding (16500, 1227)

### 1.7.12 1.5.3 Vectorizing Numerical features

#### 1.7.13 price

```
In [0]: price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset_index()
        X_train = pd.merge(X_train, price_data, on='id', how='left')
        X_cv = pd.merge(X_cv, price_data, on='id', how='left')
        X_test = pd.merge(X_test, price_data, on='id', how='left')
```

```
In [0]: from sklearn.preprocessing import Normalizer
        normalizer = Normalizer()
        # normalizer.fit(X_train['price'].values)
        # this will rise an error Expected 2D array, got 1D array instead:
        # array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
        # Reshape your data either using
        # array.reshape(-1, 1) if your data has a single feature
        # array.reshape(1, -1) if it contains a single sample.
        normalizer.fit(X_train['price'].values.reshape(-1,1))

        X_train_price_norm = normalizer.transform(X_train['price'].values.reshape(-1,1))
        X_cv_price_norm = normalizer.transform(X_cv['price'].values.reshape(-1,1))
        X_test_price_norm = normalizer.transform(X_test['price'].values.reshape(-1,1))

        print("After vectorizations")
        print(X_train_price_norm.shape, y_train.shape)
        print(X_cv_price_norm.shape, y_cv.shape)
        print(X_test_price_norm.shape, y_test.shape)
        print("=="*100)
```

After vectorizations

```
(22445, 1) (22445,)  
(11055, 1) (11055,)  
(16500, 1) (16500,)
```

### 1.7.14 tnppp

```
In [0]: from sklearn.preprocessing import Normalizer  
normalizerT = Normalizer()  
# normalizer.fit(X_train['price'].values)  
# this will rise an error Expected 2D array, got 1D array instead:  
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].  
# Reshape your data either using  
# array.reshape(-1, 1) if your data has a single feature  
# array.reshape(1, -1) if it contains a single sample.  
normalizerT.fit(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))  
  
X_train_tnppp_norm = normalizerT.transform(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))  
X_cv_tnppp_norm = normalizerT.transform(X_cv['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))  
X_test_tnppp_norm = normalizerT.transform(X_test['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))  
  
print("After vectorizations")  
print(X_train_tnppp_norm.shape, y_train.shape)  
print(X_cv_tnppp_norm.shape, y_cv.shape)  
print(X_test_tnppp_norm.shape, y_test.shape)  
print("="*100)
```

After vectorizations

```
(22445, 1) (22445,)  
(11055, 1) (11055,)  
(16500, 1) (16500,)
```

### Feature\_Aggregation

```
In [0]: feat_tnppp = list(X_train['teacher_number_of_previously_posted_projects'])  
feat_price = list(X_train['price'])  
  
In [0]: feature_agg_bow = feat_cc + feat_csc + feat_ss + feat_tp + feat_pgc + feat_bow + feat_bow_pt +  
feature_agg_tfidf = feat_cc + feat_csc + feat_ss + feat_tp + feat_pgc + feat_tfidf_e + feat_tfidf_pt
```

### 1.7.15 1.5.4 Merging all the above features

- we need to merge all the numerical vectors i.e catogorical, text, numerical vectors

## SET1

In [0]: # merge two sparse matrices: <https://stackoverflow.com/a/19710648/4084039>

```
from scipy.sparse import hstack
X_tr1 = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_grade_ohe, X_train_price_norm, X_train_project_title_bow, X_train_project_title_tfidf))
X_cr1 = hstack((X_cv_cc_ohe, X_cv_csc_ohe, X_cv_grade_ohe, X_cv_price_norm, X_cv_project_title_bow, X_cv_project_title_tfidf))
X_te1 = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_grade_ohe, X_test_price_norm, X_test_project_title_bow, X_test_project_title_tfidf))

print("Final Data matrix")
print(X_tr1.shape, y_train.shape)
print(X_cr1.shape, y_cv.shape)
print(X_te1.shape, y_test.shape)
print("=="*100)
```

Final Data matrix

```
(22445, 52054) (22445,)
(11055, 52054) (11055,)
(16500, 52054) (16500,)
```

=====

## SET2

In [0]: # merge two sparse matrices: <https://stackoverflow.com/a/19710648/4084039>

```
from scipy.sparse import hstack
X_tr2 = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_grade_ohe, X_train_price_norm, X_train_project_title_bow, X_train_project_title_tfidf))
X_cr2 = hstack((X_cv_cc_ohe, X_cv_csc_ohe, X_cv_grade_ohe, X_cv_price_norm, X_cv_project_title_bow, X_cv_project_title_tfidf))
X_te2 = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_grade_ohe, X_test_price_norm, X_test_project_title_bow, X_test_project_title_tfidf))

print("Final Data matrix")
print(X_tr2.shape, y_train.shape)
print(X_cr2.shape, y_cv.shape)
print(X_te2.shape, y_test.shape)
print("=="*100)
```

Final Data matrix

```
(22445, 10057) (22445,)
(11055, 10057) (11055,)
(16500, 10057) (16500,)
```

=====

## 2 Assignment 4: Naive Bayes

<li><strong>Apply Multinomial NaiveBayes on these feature sets</strong>

<ul>

- <li><font color='red'>Set 1</font>: categorical, numerical features + project\_title(BOW) + preprocessed\_email</li>
- <li><font color='red'>Set 2</font>: categorical, numerical features + project\_title(TFIDF)+ preprocessed\_email</li>

```

    </ul>
</li>
<br>
<li><strong>The hyper paramter tuning(find best Alpha)</strong>
    <ul>
<li>Find the best hyper parameter which will give the maximum <a href='https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/'>AUC</a> value</li>
<li>Consider a wide range of alpha values for hyperparameter tuning, start as low as 0.00001</li>
<li>Find the best hyper paramter using k-fold cross validation or simple cross validation data</li>
<li>Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter
    </ul>
</li>
<br>
<li><strong>Feature importance</strong>
    <ul>
<li>Find the top 10 features of positive class and top 10 features of negative class for both feature sets <font color='red'>learn.org/stable/modules/generated/sklearn.naive_bayes.MultinomialNB</font> and print the results
    </ul>
</li>
<br>
<li><strong>Representation of results</strong>
    <ul>
<li>You need to plot the performance of model both on train data and cross validation data for each hyper parameter. On the x-axis you will have alpha values, since they have a wide range, just to represent those alpha values on the graph, apply np.log10
<img src='train_cv_auc.JPG' width=300px></li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data
<img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpf-fpr-fnr-fnr-1/'>confusion matrix</a> with predicted and original labels
<img src='confusion_matrix.png' width=300px></li>
    </ul>
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table
    <img src='summary.JPG' width=400px>
    </li>
    </ul>
</li>

```

### Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method `fit_transform()` on your train data, and apply the method `transform()` on cv/test data.

4. For more details please go through this link.
2. Naive Bayes

## 2.1 HyperParameter Tuning

### 2.1.1 SET1

```
In [0]: # https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.GridSearchCV.html
from sklearn.model_selection import GridSearchCV
from sklearn.naive_bayes import MultinomialNB

nb1 = MultinomialNB()

parameters = {'alpha': [0.0001, 0.0025, 0.0005, 0.0075, 0.001, 0.025, .005, 0.075, 0.1, 0.25, 0.5, 0.75, 1]}
clf = GridSearchCV(nb1, parameters, cv=3, scoring='roc_auc')
mnbl = clf.fit(X_tr1, y_train)

train_auc1 = clf.cv_results_['mean_train_score']
train_auc_std1 = clf.cv_results_['std_train_score']
cv_auc1 = clf.cv_results_['mean_test_score']
cv_auc_std1 = clf.cv_results_['std_test_score']


plt.plot(parameters['alpha'], train_auc1, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['alpha'], train_auc1 - train_auc_std1, train_auc1 + train_auc_std1, alpha=0.2, color='red')

plt.plot(parameters['alpha'], cv_auc1, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['alpha'], cv_auc1 - cv_auc_std1, cv_auc1 + cv_auc_std1, alpha=0.2, color='red')

plt.scatter(parameters['alpha'], train_auc1, label='Train AUC points')
plt.scatter(parameters['alpha'], cv_auc1, label='CV AUC points')

plt.xscale('log')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

print(train_auc1)
print(cv_auc1)      #k_best = 1.0
```



```
[0.9923867 0.98895263 0.9909895 0.9869453 0.99021096 0.98373462
0.9877678 0.97920372 0.97762801 0.9709549 0.96344119 0.95754415
0.95244305]
[0.58100178 0.61471396 0.59906309 0.63055899 0.6067468 0.64522506
0.62382264 0.65615951 0.65996501 0.66686024 0.6708156 0.67218491
0.67268854]
```

## 2.1.2 SET2

```
In [0]: # https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html
from sklearn.model_selection import GridSearchCV
from sklearn.naive_bayes import MultinomialNB

nb2 = MultinomialNB()
parameters = {'alpha':[0.0001, 0.0025, 0.0005, 0.0075, 0.001, 0.025, .005, 0.075, 0.1, 0.25, 0.5, 0.75, 1]}
clf = GridSearchCV(nb2, parameters, cv=3, scoring='roc_auc')
mnb2 = clf.fit(X_tr2, y_train)

train_auc2= clf.cv_results_['mean_train_score']
train_auc_std2= clf.cv_results_['std_train_score']
cv_auc2 = clf.cv_results_['mean_test_score']
cv_auc_std2 = clf.cv_results_['std_test_score']
```



```

plt.plot(parameters['alpha'], train_auc2, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['alpha'],train_auc2 - train_auc_std2,train_auc2 + train_auc_std2,alpha=0.2,color='blue')

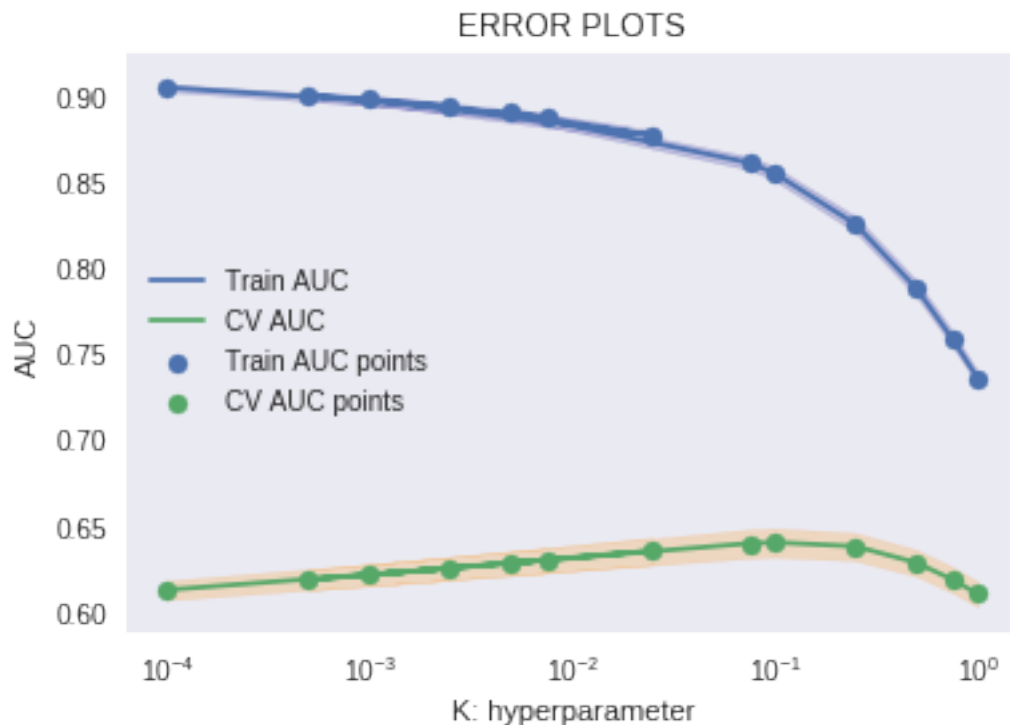
plt.plot(parameters['alpha'], cv_auc2, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['alpha'],cv_auc2 - cv_auc_std2, cv_auc2 + cv_auc_std2,alpha=0.2,color='green')

plt.scatter(parameters['alpha'], train_auc2, label='Train AUC points')
plt.scatter(parameters['alpha'], cv_auc2, label='CV AUC points')

plt.xscale('log')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

print(train_auc2)
print(cv_auc2)          #k_best = 0.1

```



[0.90473977 0.89372376 0.90025686 0.88732272 0.89776442 0.87695656

```

0.8899252 0.86094173 0.85482369 0.8253908 0.78801296 0.75906122
0.7361055 ]
[0.61292205 0.62544956 0.61889038 0.63022194 0.62166762 0.63567136
0.62845323 0.63994621 0.64053802 0.63823286 0.62874746 0.61904502
0.61071504]

```

## 2.1.3 ROC Curve

### SET1

```

In [0]: def batch_predict(clf, data):
        # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
        # not the predicted outputs

        y_data_pred = []
        tr_loop = data.shape[0] - data.shape[0]%1000
        # consider you X_tr shape is 49041, then your cr_loop will be 49041 - 49041%1000 = 49000
        # in this for loop we will iterate until the last 1000 multiplier
        for i in range(0, tr_loop, 1000):
            y_data_pred.extend(clf.predict_proba(data[i:i+1000])[:,1])
        # we will be predicting for the last data points
        y_data_pred.extend(clf.predict_proba(data[tr_loop:])[:,1])

        return y_data_pred

In [0]: # https://scikit-learn.org/https://stackoverflow.com/questions/29867367/sklearn-multinomial-nb-most-informative-features

        best_k1 = 1

        from sklearn.metrics import roc_curve, auc
        from sklearn.naive_bayes import MultinomialNB

        nb1 = MultinomialNB(alpha = best_k1)
        nb1.fit(X_tr1, y_train)
        # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
        # not the predicted outputs


        y_train_pred1 = batch_predict(nb1, X_tr1)
        y_test_pred1 = batch_predict(nb1, X_te1)

        train_fpr1, train_tpr1, tr_thresholds1 = roc_curve(y_train, y_train_pred1)
        test_fpr1, test_tpr1, te_thresholds1 = roc_curve(y_test, y_test_pred1)

        plt.plot(train_fpr1, train_tpr1, label="train AUC =" + str(auc(train_fpr1, train_tpr1)))
        plt.plot(test_fpr1, test_tpr1, label="test AUC =" + str(auc(test_fpr1, test_tpr1)))
        plt.legend()

```

```
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
```



## 2.1.4 SET2

```
In [0]: def batch_predict(clf, data):
        # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
        # not the predicted outputs

        y_data_pred = []
        tr_loop = data.shape[0] - data.shape[0]%1000
        # consider you X_tr shape is 49041, then your cr_loop will be 49041 - 49041%1000 = 49000
        # in this for loop we will iterate until the last 1000 multiplier
        for i in range(0, tr_loop, 1000):
            y_data_pred.extend(clf.predict_proba(data[i:i+1000])[:,1])
        # we will be predicting for the last data points
        y_data_pred.extend(clf.predict_proba(data[tr_loop:])[:,1])

        return y_data_pred
```

In [0]: # [https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc\\_curve.html#sklearn.metrics.roc\\_curve](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve)

```
best_k2 = 0.1
```

```
from sklearn.metrics import roc_curve, auc
from sklearn.naive_bayes import MultinomialNB
```


```
nb2 = MultinomialNB(alpha = best_k2)
nb2.fit(X_tr2, y_train)
```

```
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs
```

```
y_train_pred2 = batch_predict(nb2, X_tr2)
y_test_pred2 = batch_predict(nb2, X_te2)
```

```
train_fpr2, train_tpr2, tr_thresholds2 = roc_curve(y_train, y_train_pred2)
test_fpr2, test_tpr2, te_thresholds2 = roc_curve(y_test, y_test_pred2)
```

```
plt.plot(train_fpr2, train_tpr2, label="train AUC =" + str(auc(train_fpr2, train_tpr2)))
plt.plot(test_fpr2, test_tpr2, label="test AUC =" + str(auc(test_fpr2, test_tpr2)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
```



## 2.1.5 Confusion Matrix

### SET1

In [0]: `def predict(proba, threshold, fpr, tpr):`

```
t = threshold[np.argmax(fpr*(1-tpr))]

# (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high

print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold\n\n", np.round(t,3))
predictions = []
for i in proba:
    if i>=t:
        predictions.append(1)
    else:
        predictions.append(0)
return predictions
```

In [0]: `#https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn`

```
import seaborn as sns; sns.set()
```

```
con_m_train = confusion_matrix(y_train, predict(y_train_pred1, tr_thresholds1, train_fpr1, train_tpr1))
con_m_test = confusion_matrix(y_test, predict(y_test_pred1, te_thresholds1, test_fpr1, test_tpr1))
```

```
key = (np.asarray([['TN','FP'], ['FN', 'TP']]))
```

```
fig, ax = plt.subplots(1,2, figsize=(12,5))
```

```
labels_train = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_train.flatten())])
labels_test = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_test.flatten())])
```

```
sns.heatmap(con_m_train, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=labels_train)
sns.heatmap(con_m_test, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=labels_test)
```

```
ax[0].set_title('Train Set')
ax[1].set_title('Test Set')
```


```
plt.show()
```

the maximum value of  $tpr*(1-fpr)$  0.7073058528965523 for threshold

0.701

the maximum value of  $tpr*(1-fpr)$  0.41694824685441023 for threshold

0.998



## 2.1.6 SET2

In [0]: `def predict(proba, threshold, fpr, tpr):`

```
    t = threshold[np.argmax(fpr*(1-tpr))]
```

```
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
```

```
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold\n\n", np.round(t,3))
```

```
    predictions = []
```

```
    for i in proba:
```

```
        if i >= t:
```

```
            predictions.append(1)
```

```
        else:
```

```
            predictions.append(0)
```

```
    return predictions
```

In [0]: `#https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn`

```
import seaborn as sns; sns.set()
```

```
con_m_train = confusion_matrix(y_train, predict(y_train_pred2, tr_thresholds2, train_fpr2, train_tpr2))
```

```
con_m_test = confusion_matrix(y_test, predict(y_test_pred2, te_thresholds2, test_fpr2, test_tpr2))
```

```
key = (np.asarray([['TN','FP'], ['FN', 'TP']]))
```

```
fig, ax = plt.subplots(1,2, figsize=(12,5))
```

```

labels_train = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_
labels_test = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_

sns.heatmap(con_m_train, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], ytickl
sns.heatmap(con_m_test, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], ytickl

ax[0].set_title('Train Set')
ax[1].set_title('Test Set')

plt.show()


```

the maximum value of  $tpr*(1-fpr)$  0.541171442173297 for threshold

0.846

the maximum value of  $tpr*(1-fpr)$  0.38342887037320805 for threshold

0.868



## 2.1.7 Feature Importance

## 2.1.8 SET1

```

In [0]: pos_cls_sort_prob_1 = np.argsort((nb1.feature_log_prob_)[1])[: :-1][0:10]
neg_cls_sort_prob_1 = np.argsort((nb1.feature_log_prob_)[0])[: :-1][0:10]

```

```
pcsp = feature_agg_bow
```

```

print("Top 10 +ve class features : \n\n{0}\n\n" .format(np.take(pcsp, pos_cls_sort_prob_1)))
print("Top 10 -ve class features : \n\n{0}" .format(np.take(pcsp, neg_cls_sort_prob_1)))

```

Top 10 +ve class features :

```
['student population school' 'saying not' 'learners nannan' 'class would'
 'non english' 'laugh' 'hear daily' 'al' 'az' 'making students']
```

Top 10 -ve class features :

```
['student population school' 'saying not' 'learners nannan' 'class would'
 'non english' 'laugh' 'hear daily' 'az' 'al' 'music room']
```

### 2.1.9 SET2

```
In [0]: pos_cls_sort_prob_2 = np.argsort(abs(nb2.feature_log_prob_)[1])[:, :-1][0:10]
        neg_cls_sort_prob_2 = np.argsort(abs(nb2.feature_log_prob_)[0])[:, :-1][0:10]

        pcspt = feature_agg_tfidf

        print("Top 10 +ve class features : \n\n{0}\n\n".format(np.take(pcspt, pos_cls_sort_prob_2)))
        print("Top 10 -ve class features : \n\n{0}" .format(np.take(pcspt, neg_cls_sort_prob_2)))
```

Top 10 +ve class features :

```
['wave' 'completes' 'chooses' 'speakers' 'minds' 'reported' 'myriad'
 'speaking' 'product' 'socratic']
```

Top 10 -ve class features :

```
['redesigning' 'coloring' 'afford' 'reflect' 'coming' 'comic' 'register'
 'rehearsal' 'reliability' 'repeat']
```

## 3. Conclusions

In [0]: # Please compare all your models using Prettytable library

```
from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["Vectorizer", "Model", "HyperParameter", "AUC"]

x.add_row(["BOW", "Multinomial Naive Bayes", 1.0, 0.691])
x.add_row(["TFIDF", "Multinomial Naive Bayes", 0.1, 0.660])

print(x)
```

```
+-----+-----+-----+-----+
| Vectorizer |      Model      | HyperParameter | AUC |
```



|  |       |                         |     |       |
|--|-------|-------------------------|-----|-------|
|  |       |                         |     |       |
|  | BOW   | Multinomial Naive Bayes | 1.0 | 0.691 |
|  | TFIDF | Multinomial Naive Bayes | 0.1 | 0.66  |
|  |       |                         |     |       |

- **\*\* By seeing AUC score, we can come to conclusion that my NB model is performing average, any way better than Random model.\*\***
- **\*\* So when we take any new datapoint, there is 69% chance of getting correctly classified for SET1 and 66% chance of getting correctly classified for SET2 \*\***

### 2.1.10 Steps Followed

I took **50000 datapoints** for my analysis and building my model

- I splitted the dataset into train, cv and test dataset
- Preprocessed all the text fetatures
- Vectorized all the text, categorical and numerical features, for text i used BOW & TFIDF
- Merged all features using hstack as instructed
- Using train dataset, i plotted my AUC curve using GridSearchCV using 3Fold Cross Validation for both categories
- from AUC curve, i picked best alpha. using best alpha, i plotted ROC curve on train and test data.
- Then i plotted my confusion matrix for both the sets.
- I chose top 20 features for both +ve and -ve class for both the sets.
- Atlast you can see my result in tabular format.